

Compiling and running code on CPU nodes

Johan Hellsvik

Cray programming environment (CPE)

Reference pages: [Compilers and libraries](#), [Build systems course](#)

The Cray Programming Environment (CPE) provides consistent interface to multiple compilers and libraries.

- In practice, we recommend
 - `ml cpeCray/24.11`
 - `ml cpeGNU/24.11`
 - `ml cpeAOCC/24.11`
- The `cpeCray`, `cpeGNU` and `cpeAOCC` modules are available after `ml PDC/24.11`
- No need to `module swap` or `module unload`

Compiler wrappers

- Compiler wrappers for different programming languages
 - `cc` : C compiler wrapper
 - `cc` : C++ compiler wrapper
 - `ftn` : Fortran compiler wrapper
- The wrappers choose the required compiler version and target architecture options.
- Automatically link to MPI library and math libraries
 - MPI library: `cray-mpich`
 - Math libraries: `cray-libsci` and `cray-fftw`

Compile a simple MPI code

- `hello_world_mpi.f90`

```
program hello_world_mpi
include "mpif.h"
integer myrank,mysize,ierr
call MPI_Init(ierr)
call MPI_Comm_rank(MPI_COMM_WORLD,myrank,ierr)
call MPI_Comm_size(MPI_COMM_WORLD,mysize,ierr)
write(*,*) "Processor ",myrank," of ",mysize,": Hello World!"
call MPI_Finalize(ierr)
end program
```

```
ftn hello_world_mpi.f90 -o hello_world_mpi.x
```

What flags does the `ftn` wrapper activate?

- Use the flag `-craype-verbose`

```
ftn -craype-verbose hello_world_mpi.f90 -o hello_world_mpi.x
```

Compile a simple MPI code

```
user@uan01:> srun -n 8 ./hello_world_mpi.x
Processor      4  of      8 : Hello World!
Processor      6  of      8 : Hello World!
Processor      7  of      8 : Hello World!
Processor      0  of      8 : Hello World!
Processor      1  of      8 : Hello World!
Processor      2  of      8 : Hello World!
Processor      3  of      8 : Hello World!
Processor      5  of      8 : Hello World!
```

Compile a simple linear algebra code

[Link to the code](#)

Use cray-libsci

```
m1 PDC/24.11 cpeGNU/24.11
```

```
cc dgemm_test.c -o dgemm_test_craylibsci.x
```

Compile a simple linear algebra code

Use openblas

```
m1 PDC/24.11 openblas/0.3.29-cpeGNU-24.11
```

```
cc dgemm_test.c -o dgemm_test_openblas.x -I$EBROOTOPENBLAS/include -L$EBROOTOPENBLAS/lib -lopenblas
```

where the environment variable `EBROOTOPENBLAS` had been set when loading the OpenBLAS module. Its module file includes a statement

```
local root = "/pdc/software/24.11/eb/software/openblas/0.3.29-cpeGNU-24.11"  
setenv("EBROOTOPENBLAS", root)
```

which corresponds to an export statement

```
export EBROOTOPENBLAS=/pdc/software/24.11/eb/software/openblas/0.3.29-cpeGNU-24.11
```


Check the linked libraries

```
ldd dgemm_test_craylibsci.x
```

```
ldd dgemm_test_openblas.x
```

Check the linked libraries

```
ldd dgemm_test_craylibsci.x
```

```
...  
libsci_cray.so.6 => /opt/cray/pe/lib64/libsci_cray.so.6  
...
```

```
ldd dgemm_test_openblas.x
```

```
...  
libopenblas.so.0 => /pdc/software/24.11/eb/software/openblas/0.3.29-cpeGNU-24.11/lib/libopenblas.so.0  
...
```

Test run the dgemm_test code

- Run on a single core in the `shared` partition

```
salloc -n 1 -t 10 -p shared -A <name-of-allocation>  
srun -n 1 ./dgemm_test_craylibsci.x  
srun -n 1 ./dgemm_test_openblas.x  
exit
```

- Expected output:

2.700	4.500	6.300	8.100	9.900	11.700	13.500
4.500	8.100	11.700	15.300	18.900	22.500	26.100
6.300	11.700	17.100	22.500	27.900	33.300	38.700

Exercise: Compile and run `fftw_test` code

```
ml cray-fftw/3.3.10.9

wget https://people.math.sc.edu/Burkardt/c_src/fftw/fftw_test.c

cc --version
cc fftw_test.c -o fftw_test.x

ldd fftw_test.x

salloc -n 1 -t 10 -p shared -A <name-of-allocation>
srun -n 1 ./fftw_test.x
```

Environment variables for manual installation of software

- Environment variables for compilers

```
export CC=cc  
export CXX=CC  
export FC=ftn  
export F77=ftn
```

- Environment variables for compiler flags
 - add `-I` , `-L` , `-l` , etc. to Makefile
- Environment variables at runtime
 - prepend to `PATH` , `LD_LIBRARY_PATH` , etc.

What happens when loading a module

```
ml show elpa/2025.01.001-cpeGNU-24.11
```

```
whatis("Description: ELPA - Eigenvalue Solvers for Petaflop-Applications")
prepend_path("CMAKE_PREFIX_PATH", "/pdc/software/24.11/eb/software/elpa/2025.01.001-cpeGNU-24.11")
prepend_path("LD_LIBRARY_PATH", "/pdc/software/24.11/eb/software/elpa/2025.01.001-cpeGNU-24.11/lib")
prepend_path("LIBRARY_PATH", "/pdc/software/24.11/eb/software/elpa/2025.01.001-cpeGNU-24.11/lib")
prepend_path("PATH", "/pdc/software/24.11/eb/software/elpa/2025.01.001-cpeGNU-24.11/bin")
prepend_path("PKG_CONFIG_PATH", "/pdc/software/24.11/eb/software/elpa/2025.01.001-cpeGNU-24.11/lib/pkgconfig")
...
```

When running your own code

- Load correct programming environment (e.g. `cpeGNU`)
- Load correct dependencies (e.g. `openblas` if your code depends on it)
- Properly prepend to environment variables (e.g. `PATH` , `LD_LIBRARY_PATH`)
- Choose correct SLURM settings

Compiling and running code on GPU nodes

Johan Hellsvik

Reference pages:

[Building for AMD GPUs](#)

[Introduction to GPUs course](#)

Generalized programming for GPUs

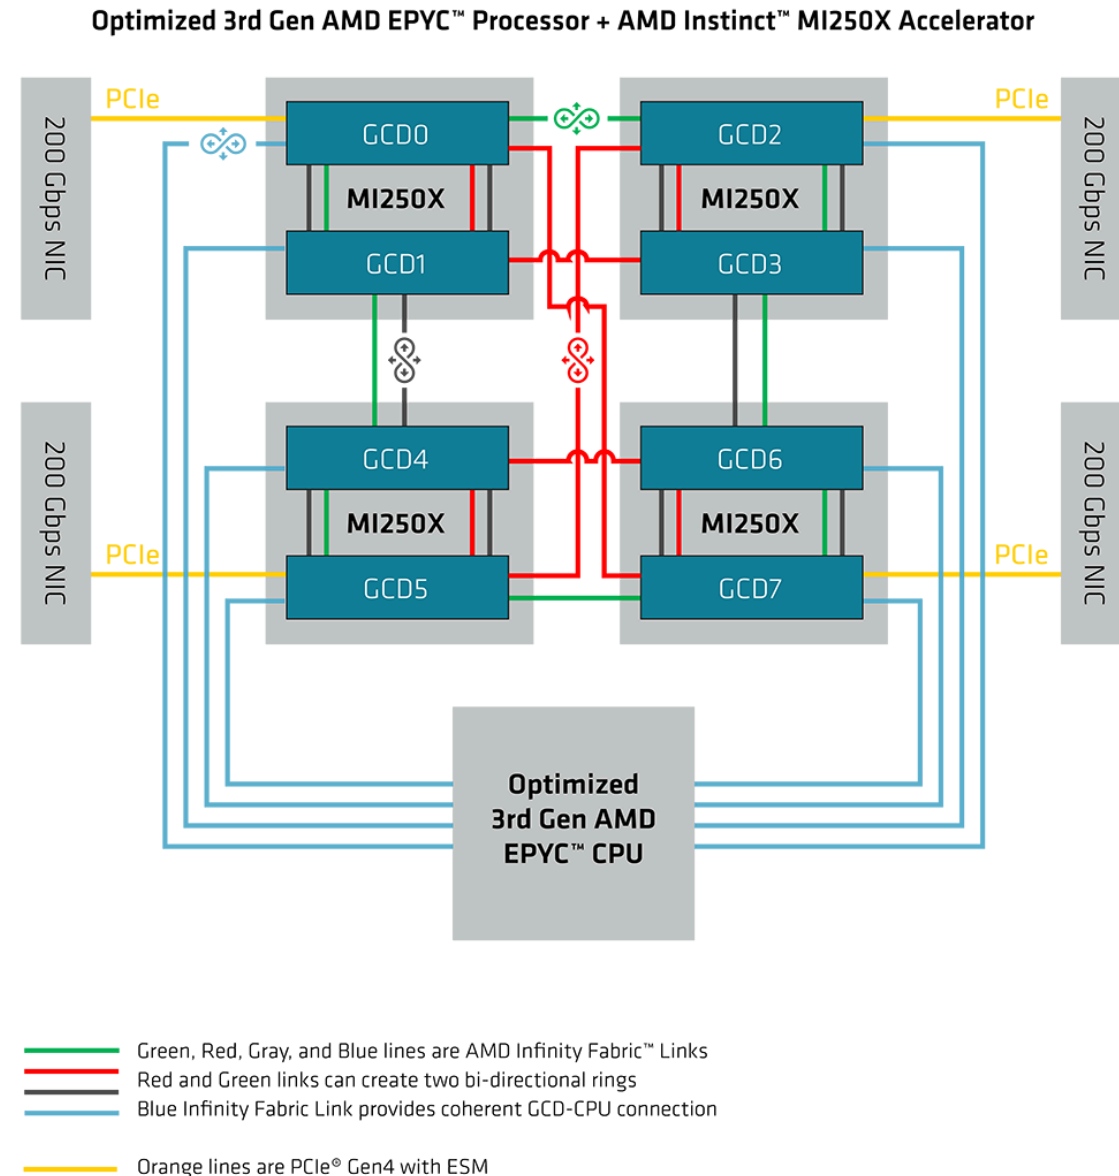
Central processing units (CPU) and graphics processing units (GPU) do different work

- CPUs have large instruction sets and execute general code.
- GPUs have smaller instructions sets. Runs compute intensive work in parallel on large number of compute units (CU).
- Code execution is started and controlled from the CPU. Compute intensive work is offloaded to the GPU.

Dardel GPU nodes

Dardel has 62 GPU nodes, each of which is equipped with

- One AMD EPYC™ processor with 64 cores
- 512 GB of shared fast HBM2E memory
- Four AMD Instinct™ MI250X GPUs (with an impressive performance of up to 95.7 TFLOPS in double precision when using special matrix operations)



AMD Radeon Open Compute (ROCm)

The AMD Radeon Open Compute (ROCm) platform is a software stack for programming and running of programs on GPUs.

- The ROCm platform supports different programming models
 - Heterogeneous interface for portability (HIP)
 - Offloading to GPU with OpenMP directives
 - The SYCL programming model
- [AMD ROCm Information Portal](#)

Setting up a GPU build environment

- Load the PDC/24.11 module and version 6.3.3 of ROCm with
 - `m1 PDC/24.11`
 - `m1 rocm/6.3.3`
- Set the accelerator target to **amd-gfx90a** (AMD MI250X GPU)
 - `m1 craype-accel-amd-gfx90a`
- Choose one of the available toolchains (Cray, Gnu, AMD)
 - `m1 cpeCray/24.11`
 - `m1 cpeGNU/24.11`
 - `m1 cpeAMD/24.11`

The ROCM info command

Information on the available GPU hardware can be displayed with the `rocm_info` command. Example output (truncated)

```
ROCK module is loaded
=====
HSA System Attributes
=====
Runtime Version:          1.1
System Timestamp Freq.:   1000.000000MHz

=====
HSA Agents
=====
*****
Agent 1
*****
      Name:                AMD EPYC 7A53 64-Core Processor
      Uuid:                 CPU-XX
```

The CRAY_ACC_DEBUG runtime environment variable

For executables that are built with the compilers of the Cray Compiler Environment (CCE), verbose runtime information can be enabled with the environment variable `CRAY_ACC_DEBUG` which takes values 1, 2 or 3.

For the highest level of information

```
export CRAY_ACC_DEBUG=3
```

Offloading to GPU with HIP

The heterogeneous interface for portability (HIP) is a hardware close (low level) programming model for GPUs. Example lines of code:

- Include statement for the HIP runtime

```
#include <hip/hip_runtime.h>
```

- HIP functions have names starting with `hip`

```
// Get number of GPUs available
if (hipGetDeviceCount(&ndevices) != hipSuccess) {
    printf("No such devices\n");
    return 1;
}
printf("You can access GPU devices: 0-%d\n", (ndevices - 1));
```

- Explicit handling of memory on the GPU

```
// Allocate memory on device
hipMalloc(&devs1, size);
hipMalloc(&devs2, size);
// Copy data host -> device
hipMemcpy(devs1, hosts1, size, hipMemcpyHostToDevice);
```

- Call to run the compute kernel on the GPU

```
// Run kernel
hipLaunchKernelGGL(MyKernel, ngrid, nblock, 0, 0, devs1, devs2);
```


Offloading to GPU with OpenMP

The OpenMP programming model can be used for directive based offloading to GPUs.

Example: A serial code that operates on arrays `vecA`, `vecB`, and `vecC`

```
! Dot product of two vectors
do i = 1, nx
    vecC(i) = vecA(i) * vecB(i)
end do
```

Implement OpenMP offloading by inserting OpenMP directives. In Fortran the directives starts with `!$omp`

```
! Dot product of two vectors
!$omp target teams distribute map(from:vecC) map(to:vecA,vecB)
do i = 1, nx
    vecC(i) = vecA(i) * vecB(i)
end do
```

Exercise 1: Hello world with HIP

Build and test run a Hello World C++ code which offloads to GPU via HIP.

- Download the [source code](https://raw.githubusercontent.com/PDC-support/introduction-to-pdc/master/example/hello_world_gpu.cpp)
 - `wget https://raw.githubusercontent.com/PDC-support/introduction-to-pdc/master/example/hello_world_gpu.cpp`
- Load the ROCm module and set the accelerator target to amd-gfx90a (AMD MI250X GPU)
 - `ml rocm/6.3.3`
 - `ml craype-accel-amd-gfx90a`
- Compile the code with the AMD hipcc compiler on the login node
 - `hipcc --offload-arch=gfx90a hello_world_gpu.cpp -o hello_world_gpu.x`

Run the code as a batch job

- Edit [job_gpu_helloworld.sh](#) to specify the compute project and reservation
- Submit the script with `sbatch job_gpu_helloworld.sh`

with program output written to `output.txt`

```
You can access GPU devices: 0-7
GPU 0: hello world``
...
```

Exercise 2: Dot product with OpenMP

Build and test run a Fortran program that calculates the dot product of vectors.

- Activate the PrgEnv-cray environment `m1 PrgEnv-cray`
- Download the [source code](https://github.com/ENCCS/openmp-gpu/raw/main/content/exercise/ex04/solution/ex04.F90)
 - `wget https://github.com/ENCCS/openmp-gpu/raw/main/content/exercise/ex04/solution/ex04.F90`
- Load the ROCm module and set the accelerator target to amd-gfx90a
 - `m1 rocm/6.3.3 craype-accel-amd-gfx90a`
- Compile the code on the login node
 - `ftn -fopenmp ex04.F90 -o ex04.x`

Run the code as a batch job

- Edit [job_gpu_ex04.sh](#) to specify the compute project and reservation
- Submit the script with `sbatch job_gpu_ex04.sh`
- with program output `The sum is: 1.25` written to `output.txt`

Optionally, test the code in interactive session.

- First queue to get one GPU node reserved for 10 minutes
 - `salloc -N 1 -t 0:10:00 -A <project name> -p gpu`
- wait for a node, then run the program `srun -n 1 ./ex04.x`
- with program output to standard out `The sum is: 1.25`

- Alternatively, login to the reserved GPU node (here nid002792) `ssh nid002792`.
- Load ROCm, activate verbose runtime information, and run the program
 - `ml rocm/6.3.3`
 - `export CRAY_ACC_DEBUG=3`
 - `./ex04.x`
- with program output to standard out

```
ACC: Version 5.0 of HIP already initialized, runtime version 50013601
ACC: Get Device 0
...
...
ACC: End transfer (to acc 0 bytes, to host 4 bytes)
ACC:
The sum is: 1.25
ACC: __tgt_unregister_lib
```